



(a) popup: 16 train, 59 holdout
holdout is 3.7x train



(b) Combined n=129.
Signal sits in the DIRTY labels.



What this corrects

Twenty notes called popup "89 rows" and treated that as the testbed size. 16 of those rows are what the model learns from.

So note 662's popup gain (+0.019) cannot come from popup's own training. It is the shared column, learned on market popup's 123 train rows and carried to popup's 59 holdout rows.

That is L2 transfer, not in-domain fit.

New rule: write domain size as train / holdout, never one number. Column cost is set by TRAIN rows; scoring is set by HOLDOUT rows.